

# Linear Regression

## Linear Algebra and its Applications

Henry Kirveslahti

AI511 University of Southern Denmark

Oct 28 2025

## Recap from yesterday

Given  $\mathbf{y} \in \mathbb{R}^n$ , design matrix  $X \in \mathbb{R}^{n \times p}$  of column rank  $p$  with  $n \geq p$ , and  $\boldsymbol{\varepsilon} \in \mathbb{R}^n$ , the linear model reads

$$\mathbf{y} = X\boldsymbol{\beta} + \boldsymbol{\varepsilon},$$

We want to find  $\boldsymbol{\beta} \in \mathbb{R}^p$  that minimizes

$$\boldsymbol{\varepsilon}^T \boldsymbol{\varepsilon} = (y - X\beta)^T(y - X\beta)$$

This is given by the least squares solution:

$$\boldsymbol{\beta} = (X^T X)^{-1} X^T \mathbf{y}.$$

Computational inconvenience: Inverting the *gram matrix*  $X^T X$ .

# Weaseling out of the inversion

Suppose we could write  $X = QR$  for some matrix  $Q$  that satisfies  $Q^T Q = I$ , and a matrix  $R$  that is in some way 'easy to invert'

The normal equation:

$$\begin{aligned} X^T y &= X^T X \beta & |X = QR \\ R^T Q^T y &= R^T Q^T Q R \beta \\ R^T Q^T y &= R^T R \beta & |\text{multiply from left with } R^{-T} \\ Q^T y &= R \beta. \end{aligned}$$

If  $R$  is easy to invert, we can recover  $\beta$  just by solving  
 $z = Q^T y = R\beta$ .

# QR Decomposition

## QR Decomposition:

$$X = QR$$

where:

- ▶  $Q \in \mathbb{R}^{n \times p}$ : matrix with orthonormal columns
- ▶  $R \in \mathbb{R}^{p \times p}$ : upper triangular matrix

## Why is this useful?

- ▶ Columns of  $Q$  form an orthonormal basis for  $\text{col}(X)$
- ▶ Makes solving least squares problems more numerically stable

## Projection Matrix via QR:

$$P = X(X^T X)^{-1} X^T = QR(R^T R)^{-1} R^T Q^T$$

$$\Rightarrow P = \boxed{QQ^T}$$

This is easy!

# QR Decomposition

- ▶ Note that  $Q^T Q = I_p$  – this is a map from  $\mathbb{R}^p$  to  $\mathbb{R}^p$
- ▶ The projection matrix  $QQ^T$  is a map from  $\mathbb{R}^n$  to  $\mathbb{R}^n$  that factors through  $\mathbb{R}^p$  - Not invertible

# QR Decomposition

- ▶ Note that  $Q^T Q = I_p$  – this is a map from  $\mathbb{R}^p$  to  $\mathbb{R}^p$
- ▶ The projection matrix  $QQ^T$  is a map from  $\mathbb{R}^n$  to  $\mathbb{R}^n$  that factors through  $\mathbb{R}^p$  - Not invertible

## Interpretation:

- ▶ The projection of  $y$  onto  $\text{col}(X)$  is:  $\hat{y} = QQ^T y$
- ▶  $QQ^T$  is the projection onto the subspace spanned by the columns of  $Q$

# Solving $R\beta = \mathbf{z}$ via Back Substitution

**Setup:**

$$R\beta = \mathbf{z}, \quad R \in \mathbb{R}^{p \times p} \text{ upper triangular, } \mathbf{z} \in \mathbb{R}^p.$$

# Solving $R\beta = \mathbf{z}$ via Back Substitution

**Setup:**

$$R\beta = \mathbf{z}, \quad R \in \mathbb{R}^{p \times p} \text{ upper triangular, } \mathbf{z} \in \mathbb{R}^p.$$

**Structure of the system:**

$$\begin{bmatrix} r_{11} & r_{12} & \cdots & r_{1p} \\ 0 & r_{22} & \cdots & r_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & r_{pp} \end{bmatrix} \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix} = \begin{bmatrix} z_1 \\ z_2 \\ \vdots \\ z_p \end{bmatrix}$$

# Solving $R\beta = \mathbf{z}$ via Back Substitution

**Setup:**

$$R\beta = \mathbf{z}, \quad R \in \mathbb{R}^{p \times p} \text{ upper triangular, } \mathbf{z} \in \mathbb{R}^p.$$

**Structure of the system:**

$$\begin{bmatrix} r_{11} & r_{12} & \cdots & r_{1p} \\ 0 & r_{22} & \cdots & r_{2p} \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & r_{pp} \end{bmatrix} \begin{bmatrix} \beta_1 \\ \beta_2 \\ \vdots \\ \beta_p \end{bmatrix} = \begin{bmatrix} z_1 \\ z_2 \\ \vdots \\ z_p \end{bmatrix}$$

**Algorithm (Back Substitution):**

$$\text{for } i = p, p-1, \dots, 1 : \quad \beta_i = \frac{1}{r_{ii}} \left( z_i - \sum_{j=i+1}^p r_{ij} \beta_j \right)$$

# Solving $R\beta = \mathbf{z}$ via Back Substitution

**Algorithm (Back Substitution):**

$$\text{for } i = p, p-1, \dots, 1 : \quad \beta_i = \frac{1}{r_{ii}} \left( z_i - \sum_{j=i+1}^p r_{ij} \beta_j \right)$$

# Solving $R\beta = \mathbf{z}$ via Back Substitution

**Algorithm (Back Substitution):**

$$\text{for } i = p, p-1, \dots, 1 : \quad \beta_i = \frac{1}{r_{ii}} \left( z_i - \sum_{j=i+1}^p r_{ij} \beta_j \right)$$

**Notes:**

- ▶ We proceed *from bottom to top*, since each equation depends only on previously computed  $\beta_j$  for  $j > i$ .
- ▶ The method requires  $r_{ii} \neq 0$  for all  $i$  (non-singular  $R$ ).
- ▶ Computational cost:  $\mathcal{O}(p^2)$ .

# Solving $R\beta = \mathbf{z}$ via Back Substitution

**Algorithm (Back Substitution):**

$$\text{for } i = p, p-1, \dots, 1 : \quad \beta_i = \frac{1}{r_{ii}} \left( z_i - \sum_{j=i+1}^p r_{ij} \beta_j \right)$$

**Notes:**

- ▶ We proceed *from bottom to top*, since each equation depends only on previously computed  $\beta_j$  for  $j > i$ .
- ▶ The method requires  $r_{ii} \neq 0$  for all  $i$  (non-singular  $R$ ).
- ▶ Computational cost:  $\mathcal{O}(p^2)$ .

**Interpretation:** Back substitution is the natural way to solve triangular systems — it is equivalent to recursively eliminating known components from the system.

# How to Compute the QR Decomposition

**Goal:** Given a full column rank matrix  $X \in \mathbb{R}^{n \times p}$ , find  $Q \in \mathbb{R}^{n \times p}$ ,  $R \in \mathbb{R}^{p \times p}$  such that:

$$X = QR$$

where:

- ▶ Columns of  $Q$  are orthonormal:  $Q^T Q = I$
- ▶  $R$  is upper triangular

**Method: Classical Gram–Schmidt Orthogonalization**

Given columns of  $X = [x_1, x_2, \dots, x_n]$ :

1. Set  $q_1 = \frac{x_1}{\|x_1\|}$
2. For  $j = 2$  to  $n$ :
  - ▶ Subtract projections onto previous  $q_i$ 's:

$$u_j = x_j - \sum_{i=1}^{j-1} \langle x_j, q_i \rangle q_i$$

- ▶ Normalize:  $q_j = \frac{u_j}{\|u_j\|}$

# QR Decomposition of an $n \times p$ Matrix

$$Q = \begin{bmatrix} | & & | \\ q_1 & \cdots & q_p \\ | & & | \end{bmatrix} \in \mathbb{R}^{n \times p}, \quad R = \begin{bmatrix} r_{11} & r_{12} & \cdots & r_{1p} \\ 0 & r_{22} & \cdots & r_{2p} \\ \vdots & \ddots & \ddots & \vdots \\ 0 & \cdots & 0 & r_{pp} \end{bmatrix} \in \mathbb{R}^{p \times p}.$$

The  $q_i$  are orthonormal vectors (basis) from the Gram-Schmidt process, and

$r_{i,j} = \langle q_i, x_j \rangle = q_i^T x_j$  (the coefficients for the basis)

# En lite visa om Gram-Schmidts Metod

There is a classic folk song on how to find an orthogonal basis [in Swedish]:

Ta en delrumsbas  $M$ , och en vektor  $a$ ;

Projisera ner, ta dess residual;

Normalisera, tillför den till  $M$ ;

Ta sen nästa vektor, börja om igen.

(The melody is that of *Itsy Bitsy Spider* [da: Lille Peter Edderkop], in case you were wondering)

# Example Computation

Let's compute the  $QR$  decomposition of

$$X = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 0 & 1 \\ 1 & 1 & 1 \\ 1 & 2 & 0 \end{pmatrix}$$

# Extending linear models

- ▶ Linear regression model needs only be linear in its parameters – but not in data!
- ▶ Given  $x, y$ , we could also just construct  $x, x^2, y$  and fit a polynomial regression  $y = \beta_0 + \beta_1 x + \beta_2 x^2$  – the powers of  $x$  are just another column
- ▶ For example, we might take log transforms of the response and or the regressors
- ▶ Recall: Exponential function  $\exp$  is a function  $f$  that satisfies  $f(x+y) = f(x)f(y)$
- ▶ Then also  $f(kx) = \underbrace{f(x)f(x)\dots f(x)}_{k \text{ times}} = f(x)^k$
- ▶ The inverse, the logarithm, turns products into sums  
 $\log(a)\log(b) = \log(a+b)$
- ▶ Consequently,  $\log(x^k) = \log(\underbrace{xx\dots x}_{k \text{ times}}) = k\log(x)$

# Always interpret on the original scale of $y$

- ▶ If we log the response, we get a model

$$\log(y) = \beta_0 + \beta_1 x$$

This is equivalent to (c.f. previous slide)

$$y = \exp(\beta_0) \exp(\beta_1 x) = a \exp(\beta_1 x)$$

- ▶ We can also take log of the regressor:

$$y = \beta_0 + \beta_1 \log(x)$$

or

$$\exp(y) = \exp(\beta_0) x^{\beta_1} = a x^{\beta_1}$$

(statistically speaking, it is usually better to interpret on the original  $y$ -scale - but this )

# Log-log transform

- ▶ We can also fit a model on log-log scale:

$$\log(y) = \beta_0 + \beta_1 \log(x)$$

This is equivalent to

$$y = \exp(\beta_0) \exp(\beta_1 \log(x)) = ax^{\beta_1}$$

# Statistical Digression: Interpreting coefficients under transformations

- ▶ Log transforming  $x$  doesn't usually need any extra interpretation (except that now the effect of  $\beta_1$  is the effect of multiplicative change in  $x$ )
- ▶ However, log-transforming the response *does affect* the interpretation of  $\beta_0 + \beta_1 x$  as the population mean – this is usually no longer accurate
- ▶ This is because if the errors  $\epsilon_i$  have certain distribution with mean 0, the mean of the exponentiated distribution is not the exponentiation of the mean the distribution
- ▶ However, exponentiation does preserve the median, so we can interpret the fitted line as the population median.
- ▶ Example: If  $y$  is for example income, our model might say something about a typical worker (i.e. the median) and not about an average worker (which is skewed by mega-millionaires)

# Core Linear Algebra Commands

Task	Command	Description
Matrix multiplication	<code>A @ B</code> or <code>np.dot(A, B)</code>	$AB$ product
Transpose	<code>A.T</code>	Transpose of $A$
Inverse	<code>np.linalg.inv(A)</code>	Inverse of square $A$
Determinant	<code>np.linalg.det(A)</code>	Determinant of $A$
Rank	<code>np.linalg.matrix_rank(A)</code>	Rank of $A$
Norm	<code>np.linalg.norm(A)</code>	Vector or matrix norm
Trace	<code>np.trace(A)</code>	Sum of diagonal elements

# Solving Linear Systems and Regression Models

Task	Command	Description
Solve $Ax = b$	<code>np.linalg.solve(A, b)</code>	Exact solution (square $A$ )
Least squares	<code>np.linalg.lstsq(A, b)</code>	Minimizes $\ Ax - b\ ^2$
Pseudo-inverse	<code>np.linalg.pinv(A)</code>	Moore–Penrose inverse
Normal equations	<code>np.linalg.inv(A.T @ A) @ A.T @ b</code>	Classic OLS formula

# Matrix Decompositions in NumPy

Decomposition	Command	Description
QR	<code>np.linalg.qr(A)</code>	$A = QR$ , $Q$ orthogonal, $R$ upper-triangular
Cholesky	<code>np.linalg.cholesky(A)</code>	$A = LL^\top$ , $A$ pos. definite
SVD	<code>np.linalg.svd(A)</code>	$A = U\Sigma V^\top$
Eigen	<code>np.linalg.eig(A)</code>	Eigenvalues/vectors of $A$
Symmetric eigendecom.	<code>np.linalg.eigh(A)</code>	Stable for symmetric $A$
Rank	<code>np.linalg.matrix_rank(A)</code>	Detects linear dependence

# Regression Workflow with NumPy

Task	Example Command	Description
Fit OLS manually	<pre>beta = np.linalg.inv(X.T @ X) @ X.T @ y</pre>	$(X^T X)^{-1} X^T y$
Fit safely	<pre>np.linalg.lstsq(X, y)</pre>	Stable least squares
Predict	<pre>y_pred = X @ beta</pre>	Compute fitted values
Residuals	<pre>res = y - y_pred</pre>	Model residuals
R-squared	<pre>1 - np.sum(res**2)/np.sum((y-y.mean())**2)</pre>	Fit quality
Back substitution	<pre>np.linalg.solve(R, z)</pre>	Solve $R\beta = z$

# Diagnostics and Stability Tools

Concept	Command	Description
Condition number	<code>np.linalg.cond(A)</code>	Numerical stability check
Check symmetry	<code>np.allclose(A, A.T)</code>	Is $A$ symmetric?
Positive definite	<code>np.all(np.linalg.eigvals(A) &gt; 0)</code>	Eigenvalues $> 0$ ?
Orthogonality check	<code>np.allclose(Q.T @ Q, I)</code>	Verifies $Q$ orthogonal
QR regression	<code>Q,R=scipy.linalg.qr(X, mode='economic')</code>	Stable OLS via QR
Solve via QR	<code>beta = scipy.linalg.solve_triangular(R, Q.T @ y)</code>	Solve $R\beta = Q^T y$

NB: These are for Scipy 1.16 - if you are using older versions your mileage may vary. Always check documentation